



Non-negative Tensor Decompositions for Unsupervised Learning

2023 SIAM PNW Conference

Western Washington University, Bellingham, WA

Nicholas Richardson

October 14, 2023



THE UNIVERSITY
OF BRITISH COLUMBIA



Table of Contents

Overview

► Overview

► Models

► Solving a model

► Special Considerations

► Conclusion



The Setting

Overview

- Do you have a lot of data?



The Setting

Overview

- Do you have a lot of data?
- A function that depends on many independent variables



The Setting

Overview

- Do you have a lot of data?
- A function that depends on many independent variables
- Many n -dimensional samples



The Setting

Overview

- Do you have a lot of data?
- A function that depends on many independent variables
- Many n -dimensional samples
- Want to compress the data



The Setting

Overview

- Do you have a lot of data?
- A function that depends on many independent variables
- Many n -dimensional samples
- Want to compress the data
- Or find relationships between samples/variables



The Setting

Overview

- Do you have a lot of data?
- A function that depends on many independent variables
- Many n -dimensional samples
- Want to compress the data
- Or find relationships between samples/variables

Tensor Decompositions are for you!



What is a Tensor Decomposition?

Overview

1. Take your data array



What is a Tensor Decomposition?

Overview

1. Take your data array
2. Select a model



What is a Tensor Decomposition?

Overview

1. Take your data array
2. Select a model
3. Solve a block-convex minimization problem between your data and a sum-product of smaller sized arrays using block coordinate gradient decent but its not converging to reasonable values because you forgot to set a unique scaling for your factors so you settle for a sub-optimal solution by projecting onto a reasonable set



What is a Tensor Decomposition?

Overview

1. Take your data array
2. Select a model
3. Write it as a product of smaller arrays



Benefits

Overview

- Less storage than original data



Benefits

Overview

- Less storage than original data
- Interpretable results



Benefits

Overview

- Less storage than original data
- Interpretable results
- Easy to implement optimization problems often converge to Nash Point



Benefits

Overview

- Less storage than original data
- Interpretable results
- Easy to implement optimization problems often converge to Nash Point
 - Minima \implies critical point \implies (local) Nash point



Benefits

Overview

- Less storage than original data
- Interpretable results
- Easy to implement optimization problems often converge to Nash Point
 - Minima \implies critical point \implies (local) Nash point
 - Some models have: critical point \iff (global) Nash point



Benefits

Overview

- Less storage than original data
- Interpretable results
- Easy to implement optimization problems often converge to Nash Point
 - Minima \implies critical point \implies (local) Nash point
 - Some models have: critical point \iff (global) Nash point
- No “training” required



Benefits

Overview

- Less storage than original data
- Interpretable results
- Easy to implement optimization problems often converge to Nash Point
 - Minima \implies critical point \implies (local) Nash point
 - Some models have: critical point \iff (global) Nash point
- No “training” required
- none or only a few hyperparameters to tune



Table of Contents

Models

► Overview

► **Models**

► Solving a model

► Special Considerations

► Conclusion



The Data Tensor

Models

- Order- N tensor

$$Y \in \mathbb{R}^{I_1 \times \cdots \times I_N}$$



The Data Tensor

Models

- Order- N tensor

$$Y \in \mathbb{R}^{I_1 \times \cdots \times I_N}$$

- Often consider additional constraints



The Data Tensor

Models

- Order- N tensor

$$Y \in \mathbb{R}^{I_1 \times \cdots \times I_N}$$

- Often consider additional constraints

- non-negative: $Y_{i_1 \dots i_N} \in \mathbb{R}_+$



The Data Tensor

Models

- Order- N tensor

$$Y \in \mathbb{R}^{I_1 \times \cdots \times I_N}$$

- Often consider additional constraints

- non-negative: $Y_{i_1 \dots i_N} \in \mathbb{R}_+$
- probability values: $0 \leq Y_{i_1 \dots i_N} \leq 1$



The Data Tensor

Models

- Order- N tensor

$$Y \in \mathbb{R}^{I_1 \times \cdots \times I_N}$$

- Often consider additional constraints

- non-negative: $Y_{i_1 \dots i_N} \in \mathbb{R}_+$
- probability values: $0 \leq Y_{i_1 \dots i_N} \leq 1$
- distributions: $1 = \sum_j Y_{i_1 \dots j \dots i_N}$



The Data Tensor

Models

- Order- N tensor

$$Y \in \mathbb{R}^{I_1 \times \cdots \times I_N}$$

- Often consider additional constraints

- non-negative: $Y_{i_1 \dots i_N} \in \mathbb{R}_+$
- probability values: $0 \leq Y_{i_1 \dots i_N} \leq 1$
- distributions: $1 = \sum_j Y_{i_1 \dots j \dots i_N}$
- binary: $Y_{i_1 \dots i_N} \in \{0, 1\}$



The Data Tensor

Models

- Order- N tensor

$$Y \in \mathbb{R}^{I_1 \times \cdots \times I_N}$$

- Often consider additional constraints

- non-negative: $Y_{i_1 \dots i_N} \in \mathbb{R}_+$
- probability values: $0 \leq Y_{i_1 \dots i_N} \leq 1$
- distributions: $1 = \sum_j Y_{i_1 \dots j \dots i_N}$
- binary: $Y_{i_1 \dots i_N} \in \{0, 1\}$


- If Y is too big to store, calculate $Y_{i_1 \dots i_N}$ on-the-fly



Tucker Decomposition

Models

- Factorize Y into a core tensor $G \in \mathbb{R}^{R_1 \times \dots \times R_N}$ and matrices $A^n \in \mathbb{R}^{I_n \times R_n}$






Tucker Decomposition

Models

- Factorize Y into a core tensor $G \in \mathbb{R}^{R_1 \times \dots \times R_N}$ and matrices $A^n \in \mathbb{R}^{I_n \times R_n}$

$$Y = G \times_1 A^1 \times_2 \dots \times_N A^N$$

$$Y_{i_1 \dots i_N} = \sum_{r_1 \dots r_N} G_{r_1 \dots r_N} A^1_{i_1 r_1} \dots A^N_{i_N r_N}$$





Tucker Decomposition


Models

- Factorize Y into a core tensor $G \in \mathbb{R}^{R_1 \times \dots \times R_N}$ and matrices $A^n \in \mathbb{R}^{I_n \times R_n}$

$$Y = G \times_1 A^1 \times_2 \dots \times_N A^N$$

$$Y_{i_1 \dots i_N} = \sum_{r_1 \dots r_N} G_{r_1 \dots r_N} A^1_{i_1 r_1} \dots A^N_{i_N r_N}$$

- Call Y a rank- (R_1, \dots, R_N) tensor, and R_n the n -rank where $R_n = \text{rank}_n(Y)$





Special Cases

Models

- Canonical polyadic decomposition (CP): $G = I$ and $R_1, \dots, R_N = R$



Special Cases

Models

- Canonical polyadic decomposition (CP): $G = I$ and $R_1, \dots, R_N = R$
- Smallest R is the tensor rank

$$Y = I \times_1 A^1 \times_2 \cdots \times_N A^N$$

$$Y_{i_1 \dots i_N} = \sum_r A_{i_1 r}^1 \cdots A_{i_N r}^N$$

e.g. low-rank decomp.: $Y = AB^\top$



Special Cases

Models

- Canonical polyadic decomposition (CP): $G = I$ and $R_1, \dots, R_N = R$
- Smallest R is the tensor rank

$$Y = I \times_1 A^1 \times_2 \cdots \times_N A^N$$

$$Y_{i_1 \dots i_N} = \sum_r A_{i_1 r}^1 \cdots A_{i_N r}^N$$

e.g. low-rank decomp.: $Y = AB^\top$

- Tucker- n : $A^{n+1}, \dots, A^N = I$ (possibly different sizes!)

$$Y = G \times_1 A^1 \times_2 \cdots \times_n A^n$$

e.g. Tucker-1: $Y = G \times_1 A$



Extensions & Additional Compression

Models

- Factorize the core symmetrically (Extend Tucker)

$$G_{r_1 r_2 r_3} = \sum_{i,j,k=1}^R B_{r_1 j k}^1 B_{i r_2 k}^2 B_{i j r_3}^3$$



Extensions & Additional Compression

Models

- Factorize the core symmetrically (Extend Tucker)

$$G_{r_1 r_2 r_3} = \sum_{i,j,k=1}^R B_{r_1 j k}^1 B_{i r_2 k}^2 B_{i j r_3}^3$$

- Tensor Trains (Extend Tucker-2)

$$Y_{i_1 \dots i_N} = \sum_{r_1, r_N} A_{i_1 r_1}^1 \underbrace{\sum_{j_2 \dots j_{N-2}} A_{r_1 i_2 j_2}^2 A_{j_2 i_3 j_3}^3 \dots A_{j_{N-2} i_{N-1} r_2}^{n-1}}_{G_{r_1 i_2 \dots i_{N-1} r_N}} A_{i_N r_N}^N$$



Extensions & Additional Compression

Models

- Factorize the core symmetrically (Extend Tucker)

$$G_{r_1 r_2 r_3} = \sum_{i,j,k=1}^R B_{r_1 j k}^1 B_{i r_2 k}^2 B_{i j r_3}^3$$

- Tensor Trains (Extend Tucker-2)

$$Y_{i_1 \dots i_N} = \sum_{r_1, r_N} A_{i_1 r_1}^1 \underbrace{\sum_{j_2 \dots j_{N-2}} A_{r_1 i_2 j_2}^2 A_{j_2 i_3 j_3}^3 \dots A_{j_{N-2} i_{N-1} r_2}^{n-1}}_{G_{r_1 i_2 \dots i_{N-1} r_N}} A_{i_N r_N}^N$$

- Factorize the matrices (Extend CP)

$$A_{i_n r}^n = \sum_k T_{i_n r k} b_k$$



Table of Contents

Solving a model

► Overview

► Models

► Solving a model

► Special Considerations

► Conclusion



The Optimization Problem

Solving a model

- May try to solve the model exactly

$$Y = f(A^1, \dots, A^N) \quad (\text{e.g. NMF: } Y = WH)$$



The Optimization Problem

Solving a model

- May try to solve the model exactly

$$Y = f(A^1, \dots, A^N) \quad (\text{e.g. NMF: } Y = WH)$$

- With noise, or imperfect modelling, better to find “closest” fit

$$\min_{A^n \in \mathcal{C}^n} D(Y, \hat{Y}) \quad \text{s.t.} \quad \hat{Y} = f(A^1, \dots, A^N)$$



The Optimization Problem

Solving a model

- May try to solve the model exactly

$$Y = f(A^1, \dots, A^N) \quad (\text{e.g. NMF: } Y = WH)$$

- With noise, or imperfect modelling, better to find “closest” fit

$$\min_{A^n \in \mathcal{C}^n} D(Y, \hat{Y}) \quad \text{s.t.} \quad \hat{Y} = f(A^1, \dots, A^N)$$

- Call the objective $D(Y, \hat{Y}) = F(A^1, \dots, A^N)$



The Optimization Problem

Solving a model

- May try to solve the model exactly

$$Y = f(A^1, \dots, A^N) \quad (\text{e.g. NMF: } Y = WH)$$

- With noise, or imperfect modelling, better to find “closest” fit

$$\min_{A^n \in \mathcal{C}^n} D(Y, \hat{Y}) \quad \text{s.t.} \quad \hat{Y} = f(A^1, \dots, A^N)$$

- Call the objective $D(Y, \hat{Y}) = F(A^1, \dots, A^N)$
 - $\|Y - \hat{Y}\|_F^2$



The Optimization Problem

Solving a model

- May try to solve the model exactly

$$Y = f(A^1, \dots, A^N) \quad (\text{e.g. NMF: } Y = WH)$$

- With noise, or imperfect modelling, better to find “closest” fit

$$\min_{A^n \in \mathcal{C}^n} D(Y, \hat{Y}) \quad \text{s.t.} \quad \hat{Y} = f(A^1, \dots, A^N)$$

- Call the objective $D(Y, \hat{Y}) = F(A^1, \dots, A^N)$
 - $\|Y - \hat{Y}\|_F^2$
 - $\sum_{i_1 \dots i_N} Y_{i_1 \dots i_N} \log \left(\frac{Y_{i_1 \dots i_N}}{\hat{Y}_{i_1 \dots i_N}} \right)$



Full Gradient Descent, Alternating Least-Squares, and Problems

Solving a model

- (Projected) Full Gradient Descent:

$$\mathcal{A} \leftarrow \arg \min_{\mathcal{B} \in \mathcal{C}} \langle \nabla F(\mathcal{A}), \mathcal{B} - \mathcal{A} \rangle + \frac{1}{2\alpha} \|\mathcal{B} - \mathcal{A}\|_F^2$$

$$(A^1, \dots, A^N) \leftarrow P_{\mathcal{C}} \left((A^1, \dots, A^N) - \alpha \nabla F(A^1, \dots, A^N) \right)$$



Full Gradient Descent, Alternating Least-Squares, and Problems

Solving a model

- (Projected) Full Gradient Descent:

$$\mathcal{A} \leftarrow \arg \min_{\mathcal{B} \in \mathcal{C}} \langle \nabla F(\mathcal{A}), \mathcal{B} - \mathcal{A} \rangle + \frac{1}{2\alpha} \|\mathcal{B} - \mathcal{A}\|_F^2$$

$$(A^1, \dots, A^N) \leftarrow P_{\mathcal{C}} \left((A^1, \dots, A^N) - \alpha \nabla F(A^1, \dots, A^N) \right)$$

- Non-expensive updates, but F not convex, not linear



Full Gradient Descent, Alternating Least-Squares, and Problems

Solving a model

- (Projected) Full Gradient Descent:

$$\mathcal{A} \leftarrow \arg \min_{\mathcal{B} \in \mathcal{C}} \langle \nabla F(\mathcal{A}), \mathcal{B} - \mathcal{A} \rangle + \frac{1}{2\alpha} \|\mathcal{B} - \mathcal{A}\|_F^2$$

$$(A^1, \dots, A^N) \leftarrow P_{\mathcal{C}} \left((A^1, \dots, A^N) - \alpha \nabla F(A^1, \dots, A^N) \right)$$

— Non-expensive updates, but F not convex, not linear

- Alternating Least-Squares:

$$A^n \leftarrow \arg \min_{A \in \mathcal{C}} \|Y - f(A^1, \dots, A^{n-1}, A, A^{n+1}, \dots, A^N)\|_F^2$$



Full Gradient Descent, Alternating Least-Squares, and Problems

Solving a model

- (Projected) Full Gradient Descent:

$$\mathcal{A} \leftarrow \arg \min_{\mathcal{B} \in \mathcal{C}} \langle \nabla F(\mathcal{A}), \mathcal{B} - \mathcal{A} \rangle + \frac{1}{2\alpha} \|\mathcal{B} - \mathcal{A}\|_F^2$$

$$(A^1, \dots, A^N) \leftarrow P_{\mathcal{C}} \left((A^1, \dots, A^N) - \alpha \nabla F(A^1, \dots, A^N) \right)$$

— Non-expensive updates, but F not convex, not linear

- Alternating Least-Squares:

$$A^n \leftarrow \arg \min_{A \in \mathcal{C}} \|Y - f(A^1, \dots, A^{n-1}, A, A^{n+1}, \dots, A^N)\|_F^2$$

— f linear in A^n , but waste time fully optimizing A^n every step



Full Gradient Descent, Alternating Least-Squares, and Problems

Solving a model

- (Projected) Full Gradient Descent:

$$\mathcal{A} \leftarrow \arg \min_{\mathcal{B} \in \mathcal{C}} \langle \nabla F(\mathcal{A}), \mathcal{B} - \mathcal{A} \rangle + \frac{1}{2\alpha} \|\mathcal{B} - \mathcal{A}\|_F^2$$

$$(A^1, \dots, A^N) \leftarrow P_{\mathcal{C}} \left((A^1, \dots, A^N) - \alpha \nabla F(A^1, \dots, A^N) \right)$$

— Non-expensive updates, but F not convex, not linear

- Alternating Least-Squares:

$$A^n \leftarrow \arg \min_{A \in \mathcal{C}} \|Y - f(A^1, \dots, A^{n-1}, A, A^{n+1}, \dots, A^N)\|_F^2$$

- f linear in A^n , but waste time fully optimizing A^n every step
- may not converge to stationary point



Block Coordinate Descent (BCD)

Solving a model

- Combine into BCD:

$$A^n \leftarrow \arg \min_{A \in \mathcal{C}} \langle \nabla_{A^n} F, A - A^n \rangle + \frac{1}{2\alpha} \|A - A^n\|_F^2$$

$$A^n \leftarrow P_{\mathcal{C}} (A^n - \alpha \nabla_{A^n} F)$$



Block Coordinate Descent (BCD)

Solving a model

- Combine into BCD:

$$A^n \leftarrow \arg \min_{A \in \mathcal{C}} \langle \nabla_{A^n} F, A - A^n \rangle + \frac{1}{2\alpha} \|A - A^n\|_F^2$$

$$A^n \leftarrow P_{\mathcal{C}} (A^n - \alpha \nabla_{A^n} F)$$

- Alternately update the factors, but don't fully optimize each factors



Block Coordinate Descent (BCD)

Solving a model

- Combine into BCD:

$$A^n \leftarrow \arg \min_{A \in \mathcal{C}} \langle \nabla_{A^n} F, A - A^n \rangle + \frac{1}{2\alpha} \|A - A^n\|_F^2$$

$$A^n \leftarrow P_{\mathcal{C}} (A^n - \alpha \nabla_{A^n} F)$$

- Alternately update the factors, but don't fully optimize each factors
- Can choose $\alpha = 1/L_n$



Block Coordinate Descent (BCD)

Solving a model

- Combine into BCD:

$$A^n \leftarrow \arg \min_{A \in \mathcal{C}} \langle \nabla_{A^n} F, A - A^n \rangle + \frac{1}{2\alpha} \|A - A^n\|_F^2$$

$$A^n \leftarrow P_{\mathcal{C}} (A^n - \alpha \nabla_{A^n} F)$$

- Alternately update the factors, but don't fully optimize each factors
- Can choose $\alpha = 1/L_n$
 - L is the Lipschitz constant of $F(A^1, \dots, A^{n-1}, \cdot, A^{n+1}, \dots, A^N)$



Block Coordinate Descent (BCD)

Solving a model

- Combine into BCD:

$$A^n \leftarrow \arg \min_{A \in \mathcal{C}} \langle \nabla_{A^n} F, A - A^n \rangle + \frac{1}{2\alpha} \|A - A^n\|_F^2$$

$$A^n \leftarrow P_{\mathcal{C}} (A^n - \alpha \nabla_{A^n} F)$$

- Alternately update the factors, but don't fully optimize each factors
- Can choose $\alpha = 1/L_n$
 - L is the Lipschitz constant of $F(A^1, \dots, A^{n-1}, \cdot, A^{n+1}, \dots, A^N)$

$f(x,y) = \text{smoothmax}(x^2+(y-2)^2, (x-2)^2+y^2)$
 $\text{smoothmax}(a,b,q) = \log_q(q^a + q^b)$
Budget of 128 iterations; $q=10^4$




Figure: Full vs Coordinate Descent



Multiplicative Updates

Solving a model

- Alternate to BCD with non-negative constraint:

$$A_{i_1 \dots i_n}^n \leftarrow A_{i_1 \dots i_n}^n \beta_{i_1 \dots i_n}^n \quad \beta \geq 0$$



Multiplicative Updates

Solving a model

- Alternate to BCD with non-negative constraint:

$$A_{i_1 \dots i_n}^n \leftarrow A_{i_1 \dots i_n}^n \beta_{i_1 \dots i_n}^n \quad \beta \geq 0$$

- Ex: $F(A, B) = \frac{1}{2} \|Y - AB\|_F^2$ use

$$\beta^1 = \frac{YB^\top}{ABB^\top} \quad \& \quad \beta^2 = \frac{A^\top Y}{A^\top AB}$$



Multiplicative Updates

Solving a model

- Alternate to BCD with non-negative constraint:

$$A_{i_1 \dots i_n}^n \leftarrow A_{i_1 \dots i_n}^n \beta_{i_1 \dots i_n}^n \quad \beta \geq 0$$

- Ex: $F(A, B) = \frac{1}{2} \|Y - AB\|_F^2$ use

$$\beta^1 = \frac{YB^\top}{ABB^\top} \quad \& \quad \beta^2 = \frac{A^\top Y}{A^\top AB}$$

- Ensures entries stay positive



Multiplicative Updates

Solving a model

- Alternate to BCD with non-negative constraint:

$$A_{i_1 \dots i_n}^n \leftarrow A_{i_1 \dots i_n}^n \beta_{i_1 \dots i_n}^n \quad \beta \geq 0$$

- Ex: $F(A, B) = \frac{1}{2} \|Y - AB\|_F^2$ use

$$\beta^1 = \frac{YB^\top}{ABB^\top} \quad \& \quad \beta^2 = \frac{A^\top Y}{A^\top AB}$$

- Ensures entries stay positive
- Equivalent to BCD with the stepsizes $\alpha_{ij}^1 = \frac{A_{ij}}{(ABB^\top)_{ij}}$ and $\alpha_{ij}^2 = \frac{B_{ij}}{(A^\top AB)_{ij}}$



Bells and Whistles

Solving a model

- Momentum:

$$\hat{A}^{(k)} = A^{(k)} + \mu^{(k)} \left(A^{(k)} - A^{(k-1)} \right)$$



Bells and Whistles

Solving a model

- Momentum:

$$\hat{A}^{(k)} = A^{(k)} + \mu^{(k)} \left(A^{(k)} - A^{(k-1)} \right)$$

- Regularizes:

$$\|A^n\|_1, \|A^n\|_2^F, \text{ etc.}$$



Bells and Whistles

Solving a model

- Momentum:

$$\hat{A}^{(k)} = A^{(k)} + \mu^{(k)} \left(A^{(k)} - A^{(k-1)} \right)$$

- Regularizes:

$$\|A^n\|_1, \|A^n\|_2^F, \text{ etc.}$$

- Line-search Stepsize:

$$\min_{\alpha > 0} F(A^1, \dots, A^{n-1}, \underbrace{A^n - \alpha \nabla_{A^n} F}_{\text{line search}}, A^{n+1}, \dots, A^N)$$



Table of Contents

Special Considerations

► Overview

► Models

► Solving a model

► **Special Considerations**

► Conclusion



Convergence

Special Considerations

- BCD converges to (global) Nash point (A^1, \dots, A^N) :

$$A^n = \arg \min_{A \in \mathcal{C}^n} F(A^1, \dots, A^{n-1}, A, A^{n+1}, \dots, A^N), \quad n = 1, \dots, N$$



Convergence

Special Considerations

- BCD converges to (global) Nash point (A^1, \dots, A^N) :

$$A^n = \arg \min_{A \in \mathcal{C}^n} F(A^1, \dots, A^{n-1}, A, A^{n+1}, \dots, A^N), \quad n = 1, \dots, N$$

- “Cannot improve objective by updating one block”



Convergence

Special Considerations

- BCD converges to (global) Nash point (A^1, \dots, A^N) :

$$A^n = \arg \min_{A \in \mathcal{C}^n} F(A^1, \dots, A^{n-1}, A, A^{n+1}, \dots, A^N), \quad n = 1, \dots, N$$

- “Cannot improve objective by updating one block”
- For block convex functions and separable constraints:



Convergence

Special Considerations

- BCD converges to (global) Nash point (A^1, \dots, A^N) :

$$A^n = \arg \min_{A \in \mathcal{C}^n} F(A^1, \dots, A^{n-1}, A, A^{n+1}, \dots, A^N), \quad n = 1, \dots, N$$

- “Cannot improve objective by updating one block”
- For block convex functions and separable constraints:

Nash point condition

Let $F(\cdot, B), F(A, \cdot)$ be convex, and $(A_0, B_0) \in \mathcal{C} = \mathcal{C}_A \times \mathcal{C}_B$ be point. Then,

$$\mathbf{0} \in \partial(F + \delta_{\mathcal{C}})(A_0, B_0) \iff \begin{array}{l} F(A_0, B_0) \leq F(A, B_0) \quad \forall A \in \mathcal{C}_A \\ F(A_0, B_0) \leq F(A_0, B) \quad \forall B \in \mathcal{C}_B \end{array} .$$



Selecting the Rank(s) R_n

Special Considerations

- Option 1: Use information about your physical system



Selecting the Rank(s) R_n

Special Considerations

- Option 1: Use information about your physical system
 - e.g. decompose piano audio into notes and amplitudes



Selecting the Rank(s) R_n

Special Considerations

- Option 1: Use information about your physical system
 - e.g. decompose piano audio into notes and amplitudes
 - $R = 88$ since there are 88 keys



Selecting the Rank(s) R_n

Special Considerations

- Option 1: Use information about your physical system
 - e.g. decompose piano audio into notes and amplitudes
 - $R = 88$ since there are 88 keys
- Option 2: When there's only 1 unknown rank;



Selecting the Rank(s) R_n

Special Considerations

- Option 1: Use information about your physical system
 - e.g. decompose piano audio into notes and amplitudes
 - $R = 88$ since there are 88 keys
- Option 2: When there's only 1 unknown rank;
 - Solve the model for all ranks $R = 1, \dots, I$


$$R = \arg \max_r \kappa(r) := \frac{F''(r)}{(1 + F'(r)^2)^{1.5}}$$



Selecting the Rank(s) R_n

Special Considerations

- Option 1: Use information about your physical system
 - e.g. decompose piano audio into notes and amplitudes
 - $R = 88$ since there are 88 keys
- Option 2: When there's only 1 unknown rank;
 - Solve the model for all ranks $R = 1, \dots, I$
 - Compute final objective value $F(R)$




$$R = \arg \max_r \kappa(r) := \frac{F''(r)}{(1 + F'(r)^2)^{1.5}}$$



Selecting the Rank(s) R_n

Special Considerations

- Option 1: Use information about your physical system
 - e.g. decompose piano audio into notes and amplitudes
 - $R = 88$ since there are 88 keys
- Option 2: When there's only 1 unknown rank;
 - Solve the model for all ranks $R = 1, \dots, I$
 - Compute final objective value $F(R)$
 - Select point of maximum curvature



$$R = \arg \max_r \kappa(r) := \frac{F''(r)}{(1 + F'(r)^2)^{1.5}}$$



Uniqueness & Scaling

Special Considerations

- These models are not (usually) unique



Uniqueness & Scaling

Special Considerations

- These models are not (usually) unique
- Ex. $Y = AB = (AC)(C^{-1}B)$ for invertible C



Uniqueness & Scaling

Special Considerations

- These models are not (usually) unique
- Ex. $Y = AB = (AC)(C^{-1}B)$ for invertable C
- Fix a scaling on factors: set $\|A^n\| = c_n$ or $\sum_j A_{...j...} = c_n$ for...



Uniqueness & Scaling

Special Considerations

- These models are not (usually) unique
- Ex. $Y = AB = (AC)(C^{-1}B)$ for invertible C
- Fix a scaling on factors: set $\|A^n\| = c_n$ or $\sum_j A_{...j...} = c_n$ for...

... all but one factor

$$\text{e.g.: } Y_{ij} = \sum_r A_{ir} B_{rj} \text{ s.t. } \|B_{r:}\|_2 = 1$$



Uniqueness & Scaling

Special Considerations

- These models are not (usually) unique
- Ex. $Y = AB = (AC)(C^{-1}B)$ for invertible C
- Fix a scaling on factors: set $\|A^n\| = c_n$ or $\sum_j A_{...j...} = c_n$ for...

... all but one factor

$$\text{e.g.: } Y_{ij} = \sum_r A_{ir} B_{rj} \text{ s.t. } \|B_{r:}\|_2 = 1$$

... all factors & add scaling parameter λ

$$\text{e.g.: } Y_{ij} = \sum_r \lambda_r A_{ir} B_{rj} \text{ s.t. } \|A_{:r}\|_2, \|B_{r:}\|_2 = 1$$



Uniqueness & Scaling

Special Considerations

- These models are not (usually) unique
- Ex. $Y = AB = (AC)(C^{-1}B)$ for invertable C
- Fix a scaling on factors: set $\|A^n\| = c_n$ or $\sum_j A_{...j...} = c_n$ for...

... all but one factor

$$\text{e.g.: } Y_{ij} = \sum_r A_{ir} B_{rj} \text{ s.t. } \|B_{r:}\|_2 = 1$$

... all factors & add scaling parameter λ

$$\text{e.g.: } Y_{ij} = \sum_r \lambda_r A_{ir} B_{rj} \text{ s.t. } \|A_{:r}\|_2, \|B_{r:}\|_2 = 1$$

- Enforce through a constraint (projection) or rescale at the end/each iteration



Uniqueness & Scaling

Special Considerations

- These models are not (usually) unique
- Ex. $Y = AB = (AC)(C^{-1}B)$ for invertable C
- Fix a scaling on factors: set $\|A^n\| = c_n$ or $\sum_j A_{...j...} = c_n$ for...

... all but one factor

$$\text{e.g.: } Y_{ij} = \sum_r A_{ir} B_{rj} \text{ s.t. } \|B_{r:}\|_2 = 1$$

... all factors & add scaling parameter λ

$$\text{e.g.: } Y_{ij} = \sum_r \lambda_r A_{ir} B_{rj} \text{ s.t. } \|A_{:r}\|_2, \|B_{r:}\|_2 = 1$$

- Enforce through a constraint (projection) or rescale at the end/each iteration
- Still only unique up to permutations of rows/columns/fibres



Table of Contents

Conclusion

► Overview

► Models

► Solving a model

► Special Considerations

► Conclusion



Summary

Conclusion

- Looked at various tensor decomposition models



Summary

Conclusion

- Looked at various tensor decomposition models
- Optimization methods to solve them



Summary







Conclusion

- Looked at various tensor decomposition models
- Optimization methods to solve them
- Practical considerations



References

Conclusion

-  W. Austin, T. G. Kolda, and T. Plantenga, "Tensor Rank Prediction via Cross Validation," Aug. 2014.
-  T. G. Kolda and B. W. Bader, "Tensor Decompositions and Applications," *SIAM Review*, vol. 51, pp. 455–500, Aug. 2009.
-  D. Lee and H. S. Seung, "Algorithms for Non-negative Matrix Factorization," in *Advances in Neural Information Processing Systems*, vol. 13, MIT Press, 2000.
-  L. Qi, Y. Chen, M. Bakshi, and X. Zhang, "Triple Decomposition and Tensor Recovery of Third Order Tensors," Mar. 2020.
-  V. Satopaa, J. Albrecht, D. Irwin, and B. Raghavan, "Finding a "Kneedle" in a Haystack: Detecting Knee Points in System Behavior," in *2011 31st International Conference on Distributed Computing Systems Workshops*, pp. 166–171, June 2011.
-  Y. Xu and W. Yin, "A Block Coordinate Descent Method for Regularized Multiconvex Optimization with Applications to Nonnegative Tensor Factorization and Completion," *SIAM Journal on Imaging Sciences*, vol. 6, pp. 1758–1789, Jan. 2013.



Thank you for listening!
Any questions?